ECE408 - Applied Parallel Programming Final Report

Neil Singh (ngsingh2) Patrick McMahon (pfmcmah2) Matthew Krikorian (krikorn2)

December 18, 2017

Milestone 1

Milestone 1 of this project was mostly setting up our test environment and preparing it to start developing the later parts of this project.

Running Time of m1.1.py

For m1.1.py, we had an accuracy of 0.8673, and an elapsed running time of 9.76 seconds.

Running Time of m1.2.py

For m1.2.py, we had an accuracy of 0.8673, and an elapsed running time of 3.19 seconds.

Most Time Consuming Kernels

Some of the most time consuming kernels when running the nvprof profile were the <code>implicit_convolve_sgemm</code> kernel, taking up 36.96% of the execution time, the <code>activation_fw_4d_kernel</code> kernel, taking up 14.33% of the execution time, and the <code>pooling_fw_4d_kernel</code> kernel, which took up 10.69% of the programs execution time.

Milestone 2

In this milestone of the project, we were tasked with writing sequential convolution code. Using the skeleton from our textbook, and the corresponding data accesses, we successfully implemented the function below.

```
{
    const int B = x.shape_[0];
    const int M = y.shape_[1];
    const int C = x.shape_[1];
    const int H = x.shape_[2];
    const int W = x.shape_[3];
    const int K = k.shape_[3];
    int H_out = H - K + 1;
    int W_out = W - K + 1;
   for (int b = 0; b < B; ++b)
     for (int m = 0; m < M; m++)</pre>
       for (int h = 0; h < H_out; h++)</pre>
         for (int w = 0; w < W_out; w++)</pre>
           y[b][m][h][w] = 0;
           for (int c = 0; c < C; c++)</pre>
             for (int p = 0; p < K; p++)</pre>
               for (int q = 0; q < K; q++)</pre>
                 y[b][m][h][w] += x[b][c][h + p][w + q] * k[m][c][p][q];
       }
}
```

Figure 1: Our CPU implementation of the convolutional layer in MxNet

Our Results For Respective Models

```
* Running nvprof python m2.1.py ece408-high 10000
New Inference
Loading fashion-mnist data... done
Loading model... done
Op Time: 8.975875
Correctness: 0.8562 Model: ece408-high
```

Figure 2: Accuracy and elapsed time for high correctness and sample size 10000

```
* Running python m2.1.py ece408-low 10000
New Inference
Loading fashion-mnist data... done
Loading model... done
Op Time: 9.003076
Correctness: 0.629 Model: ece408-low
```

Figure 3: Accuracy and elapsed time for low correctness and sample size 10000

Milestone 3

In this milestone we were tasked with writing a baseline paralllel convolution kernel that had to pass several baseline tests. This implementation is very sub-optimal, as it uses a TILE_WIDTH of 1.

```
{
   const int H_out = H - K + 1;
   const int W_out = W - K + 1;
   #define y4d(i3,i2,i1,i0) y[(i3) * (M * H_out * W_out) + (i2)*(H_out)]
        * W_out) + (i1)*(W_out) + i0]
   #define x4d(i3,i2,i1,i0) x[(i3) * (C * H * W) + (i2)*(H * W) +
        (i1)*(W) + i0]
   #define k4d(i3,i2,i1,i0) k[(i3) * (C * K * K) + (i2)*(K * K) +
        (i1)*(K) + i0]
   int W_grid = ceil(W_out / (float)TILE_WIDTH);
   int H_grid = ceil(H_out / (float)TILE_WIDTH);
   int n, m, h, w, c, p, q;
   n = blockIdx.x;
   m = blockIdx.y;
   h = blockIdx.z / W_grid + threadIdx.y;
   w = blockIdx.z % W_grid + threadIdx.x;
   float acc = 0;
   for (c = 0; c < C; c++) {</pre>
       for (p = 0; p < K; p++) {
           for (q = 0; q < K; q++) {
              if (h+p < H && w+q < W)
                  acc += x4d(n, c, h+p, w+q) * k4d(m, c, p, q);
       }
   y4d(n, m, h, w) = acc;
   #undef y4d
   #undef x4d
   #undef k4d
}
```

Figure 4: Our GPU implementation of the convolutional layer in MxNet

Nvprof GPU Profile

```
* Running nvprof python m3.1.py
Loading fashion-mnist data... done
==311== NVPROF is profiling process 311, command: python m3.1.py
Loading model... done
Op Time: 21.654475
Correctness: 0.8562 Model: ece408-high
==311== Profiling application: python m3.1.py
==311== Profiling result:
Time(%) Time Calls Avg Min Max Name
99.60% 21.6544s 1 21.6544s 21.6544s void mxnet::op::forward_kernel
```

Figure 5: Nvprof GPU profile for our forward kernel

Optimized Layer

Initial Optimization

We knew right off the bat that the first thing we were going to need to optimize was the way our implementation handled tiling. Our current implementation was giving us a slightly incorrect accuracy for any TILE_SIZE above 1, so we first had to correct that to speedup our kernel. Below is the corrected kernel from part 2, which led to an average Op time of about 450ms instead of the 14s we got with our implementation where TILE_SIZE = 1. This new implementation had TILE_SIZE = 8. This was due in part to the correction of all the floating point operations we had.

```
{
   int W_grid = (int) ceil(W_out / TILE_WIDTH * 1.0);
// int H_grid = ceil(H_out / TILE_WIDTH * 1.0);
   int n, m, h, w, c, p, q;
   n = blockIdx.x;
   m = blockIdx.y;
   h = (blockIdx.z / W_grid) * TILE_WIDTH + threadIdx.y;
   w = (blockIdx.z % W_grid) * TILE_WIDTH + threadIdx.x;
   float acc = 0.0;
   for (c = 0; c < C; c++) {
     for (p = 0; p < K; p++) {
       for (q = 0; q < K; q++) {
//
         if (h+p < H \&\& w+q < W)
           acc += x4d(n, c, h+p, w+q) * k4d(m, c, p, q);
       }
   y4d(n, m, h, w) = acc;
```

Figure 6: Our first improved GPU implementation with TILE_SIZE = 8

Further Optimization

After we made this initial optimization, we knew that a running time of about 450ms was still too poor to really say we made any real optimizations to the kernel. After this point, we chose to optimize the shared memory and tiling implementation of our kernel, to improve on the DRAM accesses our code was making. Rather than taking variables from shared memory, our implementation was making calls to global memory every time it wanted to make an operation, which is highly inefficient. To correct this, we allocated shared memory for both the inputs and the weight matrix, to allow for much faster accesses to speed up our implementation. In our new implementation, we load the filter W into shared memory, then all the threads work together to load the input into a shared memory array. Then part of the sum is computed and put into Y, and then the calculation moves onto the next channel where the threads perform the next subset of convolution. We thought this optimization would be fruitful because we were now going to use shared memory to handle our operations, which would severely decrease our running time by reducing the latency to global memory accesses. With this new and improved implementation, we were able to reduce our running time all the way down to 151ms (200ms when we run Nyprof). Below is our implementation.

Optimization Code

```
for(c = 0; c < C; c++)</pre>
      if((h0 < K) \&\& (w0 < K))
        W_{\text{shared}}[h0*K+w0] = k4d(m, c, h0, w0);
      //__syncthreads();
      int itemp = h0temp;
      for(i = h; i < h_base + X_tile_width; i += TILE_WIDTH, itemp +=</pre>
          x_tile_tile_width)
        for(j = w; j < w_base + X_tile_width; j += TILE_WIDTH)</pre>
            X_{\text{shared}}[\text{itemp} + (j - w_{\text{base}})] = x4d(n, c, i, j);
        }
      }
      __syncthreads();
      int ptemp = h0temp;
      for(p = 0; p < k2; p += K, ptemp += X_tile_width)</pre>
        for(q = 0; q < K; q++)
            acc += X_shared[ptemp + (w0 + q)] * W_shared[p+q];
        }
      }
      __syncthreads();
   }
y4d(n, m, h, w) = acc;
```

Figure 7: Our second improved GPU implementation with TILE_SIZE = 8

Nvprof GPU profile of Optimized Kernel

==317==	API calls:					
Time(%)	Time	Calls	Avq	Min	Max	Name
44.28%	1.88130s	18	104.52ms	16.944us	940.32ms	cudaStreamCreateWithFlags
27.17%	1.15452s	10	115.45ms	852ns	329.56ms	cudaFree
21.24%	902.56ms	23	39.242ms	236.93us	895.81ms	cudaMemGetInfo
4.86%	206.68ms	1	206.68ms	206.68ms	206.68ms	cudaDeviceSynchronize
1.85%	78.562ms	25	3.1425ms	5.2420us	42.421ms	cudaStreamSynchronize
0.31%	12.986ms	8	1.6233ms	17.672us	4.5404ms	cudaMemcpy2DAsync
0.15%	6.5768ms	41	160.41us	9.6770us	1.1189ms	cudaMalloc
0.03%	1.4116ms	4	352.89us	337.57us	368.23us	cuDeviceTotalMem
0.03%	1.3409ms	4	335.22us	52.578us	1.0484ms	cudaStreamCreate
0.02%	845.02us	114	7.4120us	634ns	290.39us	cudaEventCreateWithFlags
0.02%	843.11us	352	2.3950us	248ns	63.338us	cuDeviceGetAttribute
0.01%	552.32us	24	23.013us	11.030us	58.201us	cudaLaunch
0.01%	238.10us	6	39.683us	22.698us	75.956us	cudaMemcpy
0.00%	98.102us	4	24.525us	18.421us	29.756us	cuDeviceGetName
0.00%	77.039us	30	2.5670us	614ns	8.4120us	cudaSetDevice
0.00%	68.728us	104	660ns	416ns	1.9050us	cudaDeviceGetAttribute
0.00%	65.114us	145	449ns	244ns	1.3380us	cudaSetupArgument
0.00%	35.856us	2	17.928us	17.517us	18.339us	cudaStreamCreateWithPriority
0.00%	27.921us	24	1.1630us	448ns	2.8410us	cudaConfigureCall
0.00%	16.601us	10	1.6600us	1.3160us	2.1260us	cudaGetDevice
0.00%	9.2960us	17	546ns	319ns	930ns	cudaPeekAtLastError
0.00%	4.7640us	6	794ns	268ns	1.8080us	cuDeviceGetCount
0.00%	4.0080us	2	2.0040us	1.3730us	2.6350us	cudaEventRecord
0.00%	3.8340us	1	3.8340us	3.8340us	3.8340us	cudaStreamGetPriority
0.00%	3.6930us	2	1.8460us	1.5870us	2.1060us	cudaStreamWaitEvent
0.00%	3.4670us	2	1.7330us	1.5390us	1.9280us	cudaDeviceGetStreamPriorityRange
0.00%	3.3950us	6	565ns	349ns	785ns	cuDeviceGet
0.00%	3.1780us	3	1.0590us	768ns	1.2740us	cuInit
0.00%	2.5550us	5	511ns	405ns	711ns	cudaGetLastError
0.00%	2.3020us	3	767ns	532ns	1.0950us	cuDriverGetVersion
0.00%	1.3320us	1	1.3320us	1.3320us	1.3320us	cudaGetDeviceCount

Figure 8: Nvprof GPU profile for our optimized forward kernel $\,$

Summary

From the profiler, we can see a heavy increase in runtime speed from the running time of our previous implementation in milestone 3, showing that the shared memory and tiling improvements we implemented were truly worthwhile, and effective in making our forward convolution kernel much faster than it previously was. This implementation was influenced by the implementation in the textbook, which served as a great model for implementing a tiling kernel that utilized the benefits of shared memory to achieve peak performance.

Team Contributions

We all met up together and finished the sequential code for the convolutional layer in person. We also met up and wrote the parallel code together after meeting up in person, as well as crafting this report together for submission. Then for the final optimizations, we met up in person to speak about the implementation we were going to pursue, and used an online repository to assist each other in the implementation of our final kernel. We then documented our respective contributions in the report.